

Local Rule-based Explanations of Black Box Decision Systems

Riccardo Guidotti
ISTI-CNR & University of Pisa, Italy
riccardo.guidotti@isti.cnr.it

Anna Monreale
University of Pisa, Italy
anna.monreale@unipi.it

Salvatore Ruggieri
University of Pisa, Italy
salvatore.ruggieri@unipi.it

Dino Pedreschi
University of Pisa, Italy
dino.pedreschi@unipi.it

Franco Turini
University of Pisa, Italy
franco.turini@unipi.it

Fosca Giannotti
ISTI-CNR of Pisa, Italy
fosca.giannotti@isti.cnr.it

ABSTRACT

The recent years witnessed the rise of accurate but obscure decision systems which hide the logic of their internal decision processes to the users. The lack of explanations for the decisions of black box systems is a key ethical issue, and a limitation to the adoption of machine learning components in sensitive or safety-critical contexts. In this paper we focus on the problem of black box outcome explanation, i.e., explaining the reasons of the decision taken for a specific instance. We propose an agnostic method to provide interpretable and faithful explanations. **LORE** is a rule-based approach exploiting the learning of a local interpretable predictor on an artificial neighborhood generated by a genetic algorithm. **LORE** infers from the local predictor an explanation composed by a decision rule, which explains the reasons of the decision, and by a set of counterfactual rules, suggesting the changes leading to a different outcome. Wide experiments show that **LORE** outperforms existing methods and baselines both in the quality of explanations and in the accuracy in mimicking the black box.

CCS CONCEPTS

• **Information systems** → **Decision support systems**; **Data mining**; **Data analytics**;

KEYWORDS

Explanation, Decision Systems, Rules

ACM Reference Format:

Riccardo Guidotti, Anna Monreale, Salvatore Ruggieri, Dino Pedreschi, Franco Turini, and Fosca Giannotti. 1997. Local Rule-based Explanations of Black Box Decision Systems. In *Proceedings of ACM Woodstock conference (WOODSTOCK'97)*. ACM, New York, NY, USA, 10 pages. https://doi.org/10.475/123_4

1 INTRODUCTION

Popular magazines and newspapers are full of commentaries about algorithms taking critical decisions that heavily impact on our life, from granting a loan to finding a job or driving our car. The worry is not only due to the increasing automation of decision making, but mostly to the fact that the algorithms are opaque and their logic

unexplained. The main cause for this lack of transparency is that often the algorithm itself has not been directly coded by a human but it has been generated from data through machine learning. Machine learning allows building predictive models which map user features into a class (outcome or decision). This learning process is made possible by the digital records of past decisions and classification outcomes, typically provided by human experts and decision makers. The process of inferring a classification model from examples cannot be controlled step by step because the size of training data and the complexity of the learned model are too big for humans. This is how we got trapped in a paradoxical situation in which, on one side, the legislator defines new regulations requiring that automated decisions should be explained to affected people¹ while, on the other side, ever more sophisticated and obscure algorithms for decision making are generated [12, 32].

There is another aspect of learned algorithms with a potentially huge impact: they can learn bad habits from the data. If the training data contains a number of biased decision records, or misleading classification examples due to data collection mistakes or artifacts, it is likely that the resulting algorithm inherits the biases and recommends discriminatory or simply wrong decisions² [4, 5]. The inability of obtaining an explanation for what one considers a biased decision is a profound drawback of learning from big data, limiting social acceptance and trust on its adoption in many sensitive contexts. Starting from [26] a rich literature has been flourishing on discrimination discovery and avoidance. Some of the ideas developed in that context can be reinterpreted for addressing the more general problem of explaining the logic driving a decision taken by an obscure algorithm, which is precisely the problem tackled in this paper. We perform our research under some specific assumptions. First, we assume that an explanation is interesting for a user if it clarifies why a specific decision pertaining that user has been made, i.e., we aim for *local* explanations, not general, *global*, descriptions of how the system works. Second, we assume that the vehicle for offering explanations should be as close as possible to the language of reasoning, that is logic. Thus, we are also assuming that the user can understand elementary logic rules. Finally, we assume that the black box can be queried as many times as necessary, to probe its decision behavior to the scope of reconstructing its logic; this is certainly the case in a legal argumentation in court, or in an industrial setting where a company wants to stress-test a machine learning component of a manufactured product, to minimize risk of failures and consequent industrial liability. On the other hand, we make no

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).
WOODSTOCK'97, July 1997, El Paso, Texas USA
© 2016 Copyright held by the owner/author(s).
ACM ISBN 123-4567-24-567/08/06...\$15.00
https://doi.org/10.475/123_4

¹We refer here to the so-called "right to explanation" established in the European General Data Protection Regulation, entering into force in May 2018.

²www.propublica.org/article/machine-bias-risk-assessments-in-criminal-sentencing

assumptions on the specific algorithms used in the black box: we aim at an *agnostic* explanation method, one that works analyzing the input-output behavior of the black box, disregarding its internals.

We propose a solution to the black box outcome explanation problem suitable for relational, tabular data, called **LORE** (for **L**ocal **R**ule-based **E**xplanations). Given a black box binary predictor b and a specific instance x labeled with class y by b , we build a simple, interpretable predictor by first generating a balanced set of neighbor instances of the given instance x through an ad-hoc genetic algorithm, and then extracting from such a set a decision tree classifier. A *local explanation* is then extracted from the obtained decision tree. The local explanation is a pair composed by (i) a *logic rule*, corresponding to the path in the tree that explains why x has been labeled as y by b , and (ii) a set of *counterfactual rules*, explaining which conditions should be changed by x so to invert the class y assigned by b . The intuition behind our method, common to other approaches, such as LIME [27], is that the decision boundary for the black box can be arbitrarily complex over the whole data space, but in the neighborhood of a data point there is a high chance that the decision boundary is clear and simple, hence amenable to be captured by an interpretable model. The novelty of our method lies in (i) a focused procedure, based on genetic algorithm, to explore the decision boundary in the neighborhood of the data point, which produces a high-quality training data to learn the local decision tree, and (ii) a high expressiveness of the proposed local explanations, which surpasses state-of-the-art methods providing not only succinct evidence why a data point has been assigned a specific class, but also counterfactuals suggesting what should be different in the vicinity of the data point to reverse the predicted outcome. We propose extensive experiments to assess both quantitatively and qualitatively the accuracy of our explanation method.

In the rest of this paper, after describing the state of the art in the field of explanation of black box decision models (Section 2), we offer a formalization of the problem by defining the notions of *black box outcome explanation*, *explanation through interpretable models*, and *local explanation* (Section 3). We then define our method **LORE** in Section 4. Section 5 is devoted to the experiments, the set up of which requires the definition of appropriate validation measures. In the same section we critically compare local versus global explanation, and rule-based versus linear explanation. Conclusions and future research directions are discussed in Section 6.

2 RELATED WORK

Recently, the research of methods for explaining black box decision systems has caught much attention [13]. A large set of papers propose approaches for understanding the *global* logic of the black box by providing an interpretable classifier able to mimic the obscure decision system. Typically, these methods are designed for explaining specific black box models, i.e., they are not black box agnostic. Generally, the global interpretable classifier consists in a decision tree or in a set of decision rules. Decision trees have been adopted to globally explain neural networks [6, 19] and tree ensembles [14, 29]. Also, decision rules have been widely used to explain neural networks [1, 2, 17] but also to understand the global behavior of SVMs [11, 24]. Only few methods for global explanation are black box agnostic [15, 21]. In the cases in which the training

set is available, decision rules are also widely used to avoid black box models by directly designing a transparent classifier [13] which is locally or globally interpretable on its own [20, 22, 34].

Other approaches, more related to the one we propose, address the problem of explaining the *local* behavior of a black box [13]. In other words, they provide an explanation for the decision assigned to a specific record. In this context there are two kinds of approaches: the model-dependent approaches and the agnostic ones. In the first category most of the papers aim at explaining neural networks and base their explanation on saliency masks, i.e., a subset of the original record that explains what is mainly responsible for the prediction [37, 38]. Examples of salient mask are parts of an image, or words or sentences in a text. On the other hand, agnostic approaches provide explanations for any type of black box. In [27] the authors present LIME, an approach which does not depend on the type of black box. The main intuition of LIME is that the explanation can be derived locally from the instances randomly generated in the neighborhood of the instance x for which explanation should be provided, and weighted according to their proximity to it. The authors adopt linear models as comprehensible local predictor returning the importance of the features as explanation. The authors of [27] claim that LIME is also independent from the type of interpretable local classifier and from the data type. However, in practice, their algorithm and their explanation are both strictly related to the interpretable linear classifier used, i.e., a logistic regressor. Moreover, the *neighborhood generation* process adopted for images, text and tabular data is conceptually different: for images LIME randomly replaces real superpixels with superpixels containing the predominant color of the area, for text it randomly removes words, and for tabular data it assumes uniform distributions for categorical attributes and normal distributions for the continuous ones. However, a random generation of the neighborhood does not take into account and it is not able to reproduce any property of neighborhood density. Finally, since freedom from parameters indubitably promotes the usage of a methodology [18], another crucial weak point of LIME, that **LORE** overcomes, is that the number of features used in an explanation must be specified by the user.

Concerning the definition of explanation that we propose, the work [33] is related to our proposal. Indeed, [33] introduces a notion of *counterfactual* to provide an explanation. The basic idea is that an explanation consists in the smallest changes to be applied to a record x to obtain a different outcome from that one assigned by the black box to x . The counterfactuals are extracted by using an optimizer which tries to minimize the difference with x computing an x' with a different outcome. Our proposal differs from this work on the fact that in our case the counterfactual is only one of the components of an explanation. Moreover, while in [33] an optimization approach is used, we exploit a logically interpretable classifier that provides both the decision rule, describing the reason of the black box outcome, and the set of counterfactual rules, describing *what* feature to change and *how* to obtain a different decision.

To the best of our knowledge, in the literature there is no work proposing a black box agnostic method for local decision explanation based on both decision and counterfactual rules.

3 PROBLEM AND EXPLANATIONS

Let us start recalling basic notation on classification of tabular data. Afterwards, we define the *black box outcome explanation problem*, and the notion of *explanation* for which we propose a solution.

Classification, black boxes, and interpretable predictors. A *predictor* or classifier, is a function $b : \mathcal{X}^{(m)} \rightarrow \mathcal{Y}$ which maps data instances (tuples) x from a feature space $\mathcal{X}^{(m)}$ with m input features to a decision y in a target space \mathcal{Y} . We write $b(x) = y$ to denote the decision y predicted by b , and $b(X) = Y$ as a shorthand for $\{b(x) \mid x \in X\} = Y$. We restrict here to binary decisions. An instance x consists of a set of m attribute-value pairs (a_i, v_i) , where a_i is a feature (or attribute) and v_i is a value from the domain of a_i . The domain of a feature can be continuous or categorical. A predictor can be a machine learning model, a domain-expert rule-based system, or any combination of algorithmic and human knowledge processing. We assume that a predictor is available as a software function that can be queried at will. In the following, we denote by b a *black box* predictor, whose internals are either unknown to the observer or they are known but uninterpretable by humans. Examples include neural networks, SVMs, ensemble classifiers, or a composition of data mining, legacy software, and hard-coded expert systems. Instead, we denote with c an *interpretable* predictor, whose internal processing yielding a decision $c(x) = y$ can be given a symbolic interpretation understandable by a human. Examples of such predictors include rule-based classifiers, decision trees, decision sets, and rational functions.

Black Box Outcome Explanation. Given a black box predictor b and an instance x , the *black box outcome explanation problem* consists in providing an explanation e for the decision $b(x) = y$. We approach the problem by learning an interpretable predictor c that reproduces and accurately mimes the *local* behavior of the black box. An explanation of the decision is then derived from c . By *local*, we mean focusing on the behavior of the black box in the neighborhood of the specific instance x , without aiming at providing a single description of the logic of the black box for all possible instances. The neighborhood of x is not given, but rather it has to be generated as part of the explanation process. However, we assume that some knowledge is available about the characteristics of the feature space $\mathcal{X}^{(m)}$, in particular the ranges of admissible values for the domains of features and, possibly, the (empirical) distribution of features. Nothing is instead assumed about the process of constructing the black box b . Let us formalize the problem, and the approach based on interpretable models.

Definition 3.1 (Black Box Outcome Explanation). Let b be a black box, and x an instance whose decision $b(x)$ has to be explained. The *black box outcome explanation problem* consists in finding an explanation $e \in E$ belonging to a human-interpretable domain E .

Definition 3.2 (Explanation Through Interpretable Models). Let $c = \zeta(b, x)$ be an interpretable predictor derived from the black box b and the instance x using some process $\zeta(\cdot, \cdot)$. An explanation $e \in E$ is obtained through c , if $e = \epsilon(c, x)$ for some explanation logic $\epsilon(\cdot, \cdot)$ which reasons over c and x .

One point is still missing: which is a comprehensible domain E of explanations? We will define an explanation as a pair of objects.

The first component is a decision rule describing the reason for the decision value $y = c(x)$. The second component is a set of counterfactual rules, namely the minimal number of changes in the feature values of x that would reverse the decision of the predictor.

In a *decision rule* (simply, a rule) r of the form $p \rightarrow y$, the decision y is the *consequence* of the rule, while the *premise* p is a boolean condition on feature values. We assume that p is the conjunction of split conditions sc of the form³ $a \in [v_1, v_2]$, where a is a feature and v_1, v_2 are values in the domain of a plus $\pm\infty$. An instance x *satisfies* r , or r *covers* x , if the boolean condition p evaluates to true for x , i.e., if $sc(x)$ is true for every $sc \in p$. For example, the rule $\{age \leq 25, job = none\} \rightarrow deny$ is satisfied by $x_0 = \{(age=22), (job = none)\}$ and not satisfied by $x_1 = \{(age=22), (job=clerk)\}$. We say that r is *consistent* with c , if $c(x) = y$ for every instance x that satisfies r . Consistency means that the rule specifies some conditions for which the predictor makes a certain decision. When the instance x for which we have to explain the decision satisfies p , the rule $p \rightarrow y$ represents a *motivation for taking* a decision value.

Consider now a set δ of split conditions. We denote the update of p by δ as $p[\delta] = \delta \cup \{a \in [v_1, v_2] \in p \mid \nexists (a \in [w_1, w_2]) \in \delta\}$. Intuitively, $p[\delta]$ is the logical condition p with ranges for attributes overwritten as stated in δ , e.g. $\{age \leq 25, job = none\}[age > 25]$ is $\{age > 25, job = none\}$. A *counterfactual rule* for p is a rule of the form $p[\delta] \rightarrow \hat{y}$, for $\hat{y} \neq y$. If this is the case, we call δ a *counterfactual*. Consistency is meaningful also for counterfactual rules, meaning that the rule is an instance of the decision logic of c . A counterfactual δ describes *what* features to change and *how* to change them to get an outcome different from y . Since c predicts either y or \hat{y} , if such changes are applied to the given instance x , the predictor c will return a different decision. Continuing the example before, changing the age feature of x_0 to any value greater than 25 will change the predicted outcome of c from *deny* to *grant*. An expected property of a consistent counterfactual rule $p[\delta] \rightarrow \hat{y}$ is that it should be minimal w.r.t. x . Minimality is measured⁴ w.r.t. the number of split conditions in $p[\delta]$ not satisfied by x . Formally, we define $nf(p[\delta], x) = |\{sc \in p[\delta] \mid \neg sc(x)\}|$ (where $nf(\cdot, \cdot)$ stands for the number of falsified split conditions), and, when clear from the context, we simply write nf . For example, $\{age < 25, job = clerk\} \rightarrow grant$ is a counterfactual with two conditions falsified by x_0 . It is not minimal if the counterfactual $\{age > 25, job = none\} \rightarrow \hat{y}$, with only one falsified condition, is consistent for c . In summary, a counterfactual rule $p[\delta] \rightarrow \hat{y}$ is a (minimal) *motivation for reversing* a decision value.

We are now in the position to formally introduce the notion of explanation that we are able to provide.

Definition 3.3 (Local Explanation). Let x be an instance, and $c(x) = y$ be the decision of c . A local explanation $e = \langle r, \Phi \rangle$ is a pair of: a decision rule $r = (p \rightarrow y)$ consistent with c and satisfied by x ; and, a set $\Phi = \{p[\delta_1] \rightarrow \hat{y}, \dots, p[\delta_v] \rightarrow \hat{y}\}$ of counterfactual rules for p consistent with c .

³Using $\pm\infty$ we can model with a single notation typical univariate split conditions, such as equality ($a = v$ as $a \in [v, v]$), upper bounds ($a \leq v$ as $a \in [-\infty, v]$), strict lower bounds ($a > v$ as $a \in [v + \epsilon, \infty]$ for a sufficiently small ϵ).

⁴Such a measure is objective and it can be improved if additional knowledge is available on the feature domain, e.g. in order not to generate invalid or unrealistic rules. For example, which split condition between $age > 30$ and $age \leq 25$ is closer to an instance with age ? Maybe the latter is not even an option.

Algorithm 1: *LORE*(x, b)

Input : x - instance to explain, b - black box, N - # of neighbors
Output : e - explanation of x

```

1  $G \leftarrow 10$ ;  $pc \leftarrow 0.5$ ;  $pm \leftarrow 0.2$ ;           // init. parameters
2  $Z_{=} \leftarrow \text{GeneticNeigh}(x, \text{fitness}_x^x, b, N/2, G, pc, pm)$  // generate neigh.
3  $Z_{\neq} \leftarrow \text{GeneticNeigh}(x, \text{fitness}_x^x, b, N/2, G, pc, pm)$  // generate neigh.
4  $Z \leftarrow Z_{=} \cup Z_{\neq}$ ;                               // merge neighborhoods
5  $c \leftarrow \text{BuildTree}(Z)$ ;                           // build decision tree
6  $r = (p \leftarrow y) \leftarrow \text{ExtractRule}(c, x)$ ;       // extract decision rule
7  $\Phi \leftarrow \text{ExtractCounterfactuals}(c, r, x)$ ;      // extract counterfactuals
8 return  $e = \langle r, \Phi \rangle$ ;

```

This definition completes the elements of the black box outcome explanation problem. A solution to the problem will then consists of: (i) computing an interpretable predictor c for a given black box b and an instance x , i.e., defining the function $\zeta(\cdot, \cdot)$ according to Definition 3.2; (ii) deriving a local explanation from c and x , i.e., defining the explanation logic $\varepsilon(\cdot, \cdot)$ according to Definition 3.2.

4 PROPOSED METHOD

We propose **LORE** (**L**ocal **R**ule-based **E**xplanations, Algorithm 1) as a solution to the black box outcome explanation problem. An interpretable predictor c is built for a given black box b and instance x by first generating a set of N neighbor instances of x through a *genetic algorithm*, and then extracting from such a set a *decision tree* c . A local explanation, consisting of a single rule and a set of counterfactual rules, is then derived from the structure of c .

4.1 Neighborhood Generation

The goal of this phase is to identify a set of instances Z , with feature characteristics close to the ones of x , that is able to reproduce the local decision behavior of the black box b . Since the objective is to learn a predictor, the neighborhood should be flexible enough to include instances with both decision values, namely $Z = Z_{=} \cup Z_{\neq}$ where instances $z \in Z_{=}$ are such that $b(z) = b(x)$, and instances $z \in Z_{\neq}$ are such that $b(z) \neq b(x)$. In Algorithm 1, we extract balanced subsets $Z_{=}$ and Z_{\neq} (lines 2–3), and then put $Z = Z_{=} \cup Z_{\neq}$ (line 4). This task differs from approaches to instance *selection* [25], based on genetic algorithms [31] (also specialized for decision trees [36]), in that their objective is to select a subset of instances from an available training set. In our case, instead we cannot assume that the training set used to train b is available, or not even that b is a supervised machine learning predictor for which a training set exists. Our task is instead similar to instance *generation* in the field of active learning [10], also including evolutionary approaches [7]. We adopt an approach based on a *genetic algorithm* which generates $z \in Z_{=} \cup Z_{\neq}$ by maximizing the following fitness functions:

$$\text{fitness}_{=}^x(z) = I_{b(x)=b(z)} + (1 - d(x, z)) - I_{x=z}$$

$$\text{fitness}_{\neq}^x(z) = I_{b(x) \neq b(z)} + (1 - d(x, z)) - I_{x=z}$$

where $d : \mathcal{X}^{(m)} \rightarrow [0, 1]$ is a distance function, $I_{\text{true}} = 1$, and $I_{\text{false}} = 0$. The first fitness function looks for instances z similar to x (term $1 - d(x, z)$), but not equal to x (term $I_{x=z}$) for which the black box b produces the same outcome as x (term $I_{b(x)=b(z)}$). The second one leads to the generation of instances z similar to x , but

Algorithm 2: *GeneticNeigh*($x, \text{fitness}, b, N, G, pc, pm$)

Input : x - instance to explain, b - black box, fitness - fitness function, N - population size, G - # of generations, pc - crossover probability, pm - mutation probability
Output : Z - neighbors of x

```

1  $P_0 \leftarrow \{x | \forall 1 \dots N\}$ ;  $i \leftarrow 0$ ;           // population init.
2  $\text{evaluate}(P_0, \text{fitness}, b)$ ;                     // evaluate population
3 while  $i < G$  do
4    $P_{i+1} \leftarrow \text{select}(P_i)$ ;                     // select sub-population
5    $P'_{i+1} \leftarrow \text{crossover}(P_{i+1}, pc)$ ;           // mix records
6    $P''_{i+1} \leftarrow \text{mutate}(P'_{i+1}, pm)$ ;           // perform mutations
7    $\text{evaluate}(P''_{i+1}, \text{fitness}, b)$ ;               // evaluate population
8    $P_{i+1} = P''_{i+1}$ ;  $i \leftarrow i + 1$              // update population
9 end
10  $Z \leftarrow P_i$  return  $Z$ ;

```

not equal to it, for which b returns a different decision. Intuitively, for an instance z_0 such that $b(x) \neq b(z_0)$ and $x \neq z_0$, it turns out $\text{fitness}_{=}^x(z_0) < 1$. For any instance z_0 such that $b(x) = b(z_0)$, instead, we have $\text{fitness}_{=}^x(z_0) \geq 1$. Finally, $\text{fitness}_{=}^x(x) = 1$. Thus, maximization of $\text{fitness}_{=}^x(\cdot)$ occurs necessarily for instances different from x and whose prediction is equal to $b(x)$.

Like neural networks, genetic algorithms [16] are based on the biological metaphor of evolution. They have three distinct aspects. (i) The potential solutions of the problem are encoded into representations that support the variation and selection operations. In our case these representations, generally called chromosomes, correspond to instances in the feature space \mathcal{X}^m . (ii) A fitness function evaluates which chromosomes are the “best life forms”, that is, most appropriate for the result. These are then favored in *survival* and *reproduction*, thus shaping the next generation according to the fitness function. In our case, these instances correspond to those similar to x , according to distance $d(\cdot, \cdot)$, and with the same/different outcome returned by the black box b depending on the fitness function $\text{fitness}_{=}^x$ or fitness_{\neq}^x . (iii) Mating (called crossover) and mutation produce a new generation of chromosomes by recombining features of their parents. The final generation of chromosomes, according to a stopping criterion, is the one that most fit the solution.

Algorithm 2 generates the neighborhoods $Z_{=}$ and Z_{\neq} of x by instantiating the evolutionary approach of [3]. Using the terminology of [7], it is an instance of generational genetic algorithms for evolutionary prototype generation. However, while prototypes are a condensed subset of a training set that optimized performance of a predictor, we aim at generating new instances that separate well the decision boundary of the black box b . However, the classifier they use is always the one for which the population must be selected or generated and not another one like in our case. The algorithm first initializes the population P_0 with N copies of the instance x to explain. Then, it enters the evolution loop that begins by *selection* of the P_{i+1} population having the highest fitness score. After that, the crossover operator is applied on a proportion of P_{i+1} according to the pc probability, the resulting and the untouched individuals are placed in P'_{i+1} . We use a *two-point crossover* which selects two parents and two crossover features at random, and then swap the crossover feature values of the parents. Thereafter, a proportion




Figure 1: Random forest black box: purple vs green decision. Starred instance x . (Top) Uniformly random (left) and genetic generation (right). (Bottom) Density of random (left) and genetic (right) generation. (best view in color).

of P'_{i+1} , determined by pm , is mutated and placed in P''_{i+1} . The unmutated individuals are also added to P''_{i+1} . Mutation consists of replacing features values at random according to the empirical distribution⁵ of a feature. Individuals in P''_{i+1} are evaluated according to the fitness function, and the evolution loop continues until G generations are completed. Finally, the best individuals according to the fitness function are returned.

Figure 1 shows an example of neighborhood generation for a black box consisting of a random forest model and a bi-dimensional feature space. The figure contrasts uniform random generation around a specific instance x (starred) to our genetic approach. The latter yields a neighborhood that is denser in the boundary region of the predictor. The density of generated instances will be a key factor in extracting local (interpretable) predictors.

Distance Function. A key element in the definition of the fitness functions is the distance $d(x, z)$. We account for the presence of mixed types of features by a weighted sum of simple matching coefficient for categorical features, and of the normalized Euclidean distance⁶ for continuous features. Formally, assuming h categorical features and $m - h$ continuous ones, we use:

$$d(x, z) = \frac{h}{m} \cdot \text{SimpleMatch}(x, z) + \frac{m - h}{m} \cdot \text{NormEuclid}(x, z).$$

Our approach is parametric to $d(\cdot, \cdot)$, and it can readily be applied to improved heterogeneous distance functions [23, 35]. However, while more complex distances can overcome known problems such as overweighting correlated features, we point out that this is not the case when feature values are generated (at random) independently from each other, as in our approach.

4.2 Local Rule-Based Classifier and Explanation Extraction

Given the neighborhood Z of x , the second step is to build an interpretable predictor c trained on the instances $z \in Z$ labeled

⁵In experiments, we derive such a distribution from the test set of cases to explain.

⁶reference.wolfram.com/language/ref/NormalsquaredEuclideanDistance.html

Algorithm 3: ExtractCounterfactuals(c, r, x)

Input : c - decision tree, r - rule ($p \rightarrow y$), x - instance to explain
Output : Φ - set of counterfactual rules for p

```

1  $Q \leftarrow \text{GetPathsWithDifferentLabel}(c, y);$  // get paths with label  $\hat{y} \neq y$ 
2  $\Phi \leftarrow \emptyset; \min \leftarrow +\infty;$  // init. counterfactual set
3 for  $q \in Q$  do
4    $q_{len} \leftarrow nf(q, x) = |\{sc \in q \mid \neg sc(x)\}|$ 
5   if  $q_{len} < \min$  then  $\Phi \leftarrow \{q \rightarrow \hat{y}\}; \min \leftarrow q_{len};$ 
6   else if  $q_{len} = \min$  then  $\Phi \leftarrow \Phi \cup \{q \rightarrow \hat{y}\};$ 
7 end
8 return  $\Phi;$ 

```




Figure 2: Example decision tree.

with the black box decision $b(z)$. Such a predictor is intended to mimic the behavior of b locally in the Z neighborhood. Also, c must be interpretable, so that an explanation for x (decision rule and counterfactuals) can be extracted from it. The **LORE** approach considers decision tree classifiers due to the following reasons: (i) decision rules can naturally be derived from a root-leaf path in a decision tree; and, (ii) counterfactuals can be extracted by symbolic reasoning over a decision tree. For a decision tree c , we derive an explanation $e = \langle r, \Phi \rangle$ as shown in Algorithm 3. The decision rule $r = (p \rightarrow y)$ is formed by including in p the split conditions on the path from the root to the leaf node that is satisfied by the instance x , and setting $y = c(x)$. By construction, the rule r is consistent with c and satisfied by x . Consider now the counterfactual rules in Φ . The idea is to look for all paths in the decision tree c leading to a decision $\hat{y} \neq y$. Fix one of such paths, and let q be the conjunction of split conditions in it. Again by construction, $q \rightarrow \hat{y}$ is a counterfactual rule consistent with c . The counterfactual δ for which $q = p[\delta]$ has not to be explicitly computed⁷ – this is a benefit of using decision trees. Among all such q 's, only the ones with minimum number of split conditions sc not satisfied by x (line 4 of Algorithm 3) are kept in Φ . As an example, consider the decision tree in Figure 2, and the instance $x = \{(age, 22), (job, clerk), (income, 800)\}$ for which the decision *deny* (e.g., of a loan) has to be explained. The path followed by x is the leftmost one in the tree. The decision rule extracted from the path is $\{age \leq 25, job = clerk, income \leq 900\} \rightarrow deny$. There are four paths leading to the opposite decision: $q_1 = \{age \leq 25, job = clerk, income > 900\}$, $q_2 = \{17 < age \leq 25, job = other\}$, $q_3 = \{age > 25, income \leq 1500, job = other\}$, $q_4 = \{age > 25, income > 1500\}$. It turns out: $nf(q_1, x) = 1$, $nf(q_2, x) = 1$, $nf(q_3, x) = 2$, and $nf(q_4, x) = 2$. Thus, $\Phi = \{q_1 \rightarrow deny, q_2 \rightarrow deny\}$.

⁷It can be done as follows. Consider the path from the leaf of p to the leaf of q . When moving from a child to a father node, we retract the split condition. E.g., $a \leq v_2$ is retracted from $\{a \in [v_1, v_2]\}$ by adding $a \in [v_1, +\infty]$ to δ . When moving from a father node to a child, we add the split condition to δ .

As a further output **LORE** computes a *counterfactual instance* starting from a counterfactual rule $q \rightarrow \hat{y}$ and x . Among all possible instances that satisfy q , we choose the one that minimally changes attributes from x according to q . This is done by looking at the split conditions falsified by x : $\{sc \in q \mid \neg sc(x)\}$, and modifying the lower/upper bound in sc that is closer to the value in x . As an example, for the above path q_1 , the counterfactual instance of x is $\{(age, 22), (job, clerk), (income, 900 + \epsilon)\}$, and for q_2 is $\{(age, 22), (job, other), (income, 800)\}$.

5 EXPERIMENTS

LORE has been developed in Python⁸, using, for the genetic neighborhood generation, the *deap* library [9], and for decision tree induction (our interpretable predictor), the *yadt* system [28], which is a C4.5 implementation with multi-way splits of categorical attributes. In this section, after presenting the experimental setup we report (i) some analyses on the effect of the genetic algorithm parameters for the neighborhood generation; (ii) evidence that the local genetic neighborhood is more effective than a global approach and than existing oversampling and instance selection techniques; (iii) a qualitative and quantitative comparison with naïve baselines and state of the art competitors.

5.1 Experimental Setup

We ran experiments on three real-world *tabular* datasets: *adult*, *compas* and *german*⁹. In each of them, an instance represents attributes of an individual person. Explaining decisions taken by black boxes for such instances is a relevant problem. All datasets includes both categorical and continuous features.

The *adult* dataset from UCI Machine Learning Repository, includes 48,842 instances with demographic information like age, workclass, marital-status, race, capital-loss, capital-gain etc. The income divides the population into two classes “ $\leq 50K$ ” and “ $> 50K$ ”.

The *compas* dataset from ProPublica contains the features used by the COMPAS algorithm for scoring defendants and their risk (Low, Medium and High), for over 10,000 individuals. We considered two classes “Low-Medium” and “High” risk, and we use the following features: age, sex, race, priors_count, days_b_screening_arrest, length_of_stay, c_charge_degree, is_recid.

In the *german* dataset from UCI Machine Learning Repository each person of the 1,000 entries is classified as a “good” or “bad” creditor according to attributes like age, sex, checking_account, credit_amount, duration, purpose, etc.

We experimented the following predictors as black boxes: support vector machines with RBF kernel (**SVM**), random forests with 100 trees (**RF**), and multi-layer neural networks with ‘lbfgs’ solver (**NN**). Implementations of the predictors are from the *scikit-learn* library. Unless differently stated, default parameters were used. Missing values were replaced by the mean for continuous features and by the mode for categorical ones.

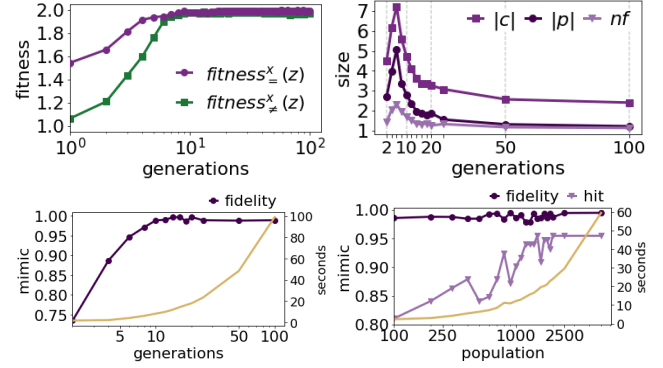


Figure 3: Impact of the number of generations G and of population size N parameters of the genetic neighborhood generation. Bottom plots also report elapsed running times.

Each dataset was randomly split into train (80% instances), and test (20% instances). The former is used to train black box predictors. The latter, denoted by X , is the set of instances for which the predicted decision needs to be explained. In the following, for some fixed set of instances, we denote by Y the set of decisions provided by the interpretable predictor c and by \hat{Y} the decisions provided by the black box b on the same set.

We consider the following properties in evaluating the mimic performances of the decision tree c inferred by **LORE** and of the explanations returned by it against the black box classifier b :

- **fidelity**(Y, \hat{Y}) $\in [0, 1]$. It compares the predictions of c and of the black box b on the instances Z used to train c [8]. It answer the question: how good is c at mimicking b ?
- **l-fidelity**(Y, \hat{Y}) $\in [0, 1]$. It compares the predictions of c and b on the instances Z_x covered by the decision rule in a local explanation for x . It answer the question: how good is the decision rule at mimicking b ?
- **cl-fidelity**(Y, \hat{Y}) $\in [0, 1]$. It compares the predictions of c and b on the instances Z_x covered by the counterfactual rules in a local explanation for x .
- **hit**(y, \hat{y}) $\in \{0, 1\}$. It compares the predictions of c and b on the instance x under analysis. It returns 1 if $y = c(x)$ is equal to $\hat{y} = b(x)$, and 0 otherwise.
- **c-hit**(y, \hat{y}) $\in \{0, 1\}$. It compares the predictions of c and b on a counterfactual instance of x built from counterfactual rules in a local explanation of x .

We measure the first three of them by the f1-measure [30]. Aggregated values of f1 and hit/c-hit are reported by averaging them over the the set of test instances $x \in X$.

5.2 Analysis of Neighborhood Generation

We analyze here the impact of the number of generations G and size of neighborhood N on the performances of instance generation and on the size complexity of the **LORE** output. We report only results for *german* dataset, since we get similar results for the other ones. The other parameters of Algorithm 2 (probabilities of crossover pc and mutation pm) are set with the default values of 0.5 and 0.2 respectively [3]. Figure 3 shows in the top plots the value of fitness functions and measures of sizes of local classifier c (decision tree

⁸The Python code and dataset for replicating the experiments will be available at <https://github.com/riccotti/LORE>. The experiments were performed on Ubuntu 16.04.1 LTS 64 bit, 32 GB RAM, 3.30GHz Intel Core i7

⁹<https://archive.ics.uci.edu/ml/datasets/adult>, <https://github.com/propublica/compas-analysis>, [https://archive.ics.uci.edu/ml/datasets/statlog+\(german+credit+data\)](https://archive.ics.uci.edu/ml/datasets/statlog+(german+credit+data))

Dataset	Method	hit	fidelity	l-fidelity	tree depth
adult	lore	.912 ± .29	.959 ± .17	.892 ± .29	4.16 ± 0.21
	global	.901 ± .28	.750 ± .00	.873 ± .27	12.00 ± 0.00
compas	lore	.942 ± .23	.992 ± .03	.937 ± .23	4.72 ± 2.15
	global	.902 ± .29	.935 ± .00	.857 ± .29	12.00 ± 0.00
german	lore	.925 ± .26	.988 ± .07	.920 ± .26	4.95 ± 2.54
	global	.880 ± .32	.571 ± .00	.824 ± .31	6.00 ± 0.00

Table 1: Local vs global approach.

depth), of decision rule (size of the antecedent p), and of counterfactual rules (number nf of falsified split conditions). The bottom plots show fidelity (f1-measure) and hit (rate) as well as running times of neighborhood generation. Fixed $N = 1000$, after 10 generations, the fitness function converges around the optimal value (top left), fidelity is almost maximized (bottom left), and also the measures of sizes (top right) become stable and small. We then set $G=10$ in all other experiments. Figure 3-(bottom right) shows instead that the size N of the neighborhood instances to be generated is relevant for the *hit* rate but not for *fidelity*. By taking into account also the running time (right side scale of the bottom plots), a good trade-off is obtained by setting $N=1000$.

5.3 Validation of Local Explanations

In this section we show the performance of our method for the local explanation extraction comparing it with a global approach and different solutions for local instance generations.

Local vs Global Explanations. Extracting a predictor from the neighborhood of an instance is a winning strategy, if contrasted to an approach that builds a single predictor from all instances in the test set, i.e., $Z = X$. In particular, this means that the interpretable predictor will be the same for all instances in the test set. Let compare our approach (reported as **LORE** in tables) with such a global approach. Table 1 reports the mean and standard deviations values of *hit*, *fidelity*, *fairness* and *tree depth* for each dataset aggregating over the results of the various black boxes. While for *hit* both **LORE** and **global** obtain similar high performances, for the other scores **LORE** considerably overtakes **global**. In particular, the size and depth of the decision tree of the global approach may lead to explanations (decision rules and counterfactuals) more complex to understand than in the local approach.

Comparing Neighborhood Generations. After concluding that “local is better than global”, we now show that our genetic programming approach improves over the following baselines in the generation of neighborhoods:

- **crn** returns as Z the $k = 10$ instances from X (the test set) that are *closest* to x ;
- **rnd** augment the output of **crn** with additional randomly generated instances so that a stratified Z is obtained;
- **ris** starting from the output of **rnd** performs the instance selection procedure¹⁰ **CNN** [25];
- **ros** starting from the output of **rnd** performs a random oversampling to balance the decision outcomes in Z .

¹⁰<http://contrib.scikit-learn.org/imbalanced-learn/>

Method	hit	fidelity	l-fidelity	c-hit	cl-fidelity
lore	.962 ± .19	.993 ± .04	.959 ± .19	.588 ± .42	.756 ± .40
crd	.924 ± .26	.855 ± .23	.894 ± .25	.349 ± .26	.583 ± .48
rnd	.946 ± .22	.904 ± .15	.920 ± .22	.494 ± .24	.712 ± .40
ris	.916 ± .27	.869 ± .05	.870 ± .26	.501 ± .22	.708 ± .39
ros	.968 ± .17	.965 ± .03	.953 ± .17	.491 ± .22	.733 ± .34

Table 2: Comparison of neighborhood generations methods.

Dataset	german		compass		adult	
Black Box	lore	lime	lore	lime	lore	lime
RF	.925 ± .2	.880 ± .3	.941 ± .2	.826 ± .4	.901 ± .3	.824 ± .4
NN	.980 ± .1	1.00 ± .0	.987 ± .1	.902 ± .3	.918 ± .3	.998 ± .1
SVM	1.00 ± .0	.966 ± .1	.997 ± .1	.900 ± .3	.985 ± .1	.987 ± .1

Table 3: LORE vs LIME: hit scores.

Table 2 reports the aggregated evaluation measures over the various black boxes and datasets. **LORE** overtook the performance of all the other neighbors generators. Intuitively, this means that **LORE** contributes more than the other methods in capturing/explaining the behavior of the black box, both for direct and counterfactual decisions. Such a conclusion is reinforced by Figure 4, which shows the box plots of the distributions of *fidelity*, *l-fidelity* and *cl-fidelity*, and some summary data on the size of decision trees ($|c|$), of decision rule premises ($|p|$), and of the number of falsified split conditions in counterfactual rules (nf). **LORE** has the highest mean and median f1-measures (high mimic of the black box), the smallest interquartile ranges (low variability of results), and the lowest complexity sizes.

5.4 Rule-Based vs Linear Explanations

In this section we compare our notion of explanation with some approaches from the state of the art.

5.4.1 Rules vs linear regression for explanations. We present first a quantitative and qualitative comparison with the linear explanations of LIME¹¹ [27]. A first crucial difference is that in LIME, the number of features composing an explanation is an input parameter that must be specified by the user. **LORE**, instead, automatically provides the user with an explanation including only the features useful to justify the black box decision. This is a clear improvement over LIME. In experiments, unless otherwise stated, we vary the number of features of LIME explanations from two to ten and we consider the performance with the highest score.

Quantitative Comparison. Table 3 reports the mean and standard deviation of *hit* for each black box predictor and dataset. Moreover, Figure 5 details the box plots of *fidelity* (top) and *l-fidelity* (bottom). Results show that **LORE** definitely outperforms LIME under various viewpoints. Regarding the *hit* score, even when **LORE** is worse than LIME, it has a score close to 1. For RF black box, instead, LIME performs considerably worse than **LORE**. The box plots show that, in addition, **LORE** has better (local) fidelity scores and is more robust than LIME, which, on the contrary, exhibits very high variability in the neighborhood of the instance to explain (i.e., for *l-fidelity*). This can be tracked back to the genetic instance

¹¹<https://github.com/marcotcr/lime>




Figure 4: Comparison of neighborhood generations methods.




Figure 5: LORE vs LIME: box plots of fidelity and l-fidelity. Numbers on top are the mean values.




Figure 6: Neighborhoods of LORE (left) and LIME (right).

generation of **LORE**. Figure 6 reports a multidimensional scaling of the neighborhood of a sample instance x generated by the two approaches. **LORE** computes a dense and compact neighborhood. The instances generated by **LIME**, instead, can be very distant from each other and always with a low density around x .

Qualitative Comparison. We claim that the explanation provided by **LORE** in terms of decision and counterfactual rules is more abstract and comprehensible than the one of **LIME**. Consider the example in Figure 7. The top part reports a **LORE** local explanation for an instance x from the german dataset. The bottom part a **LIME** explanation. Weights are associated to the categorical values in the instance x to explain, and to continuous upper/lower bounds where the bounding values are taken from x . Each weight tells the user how much the decision would have changed for different (resp., smaller/greater) values of a specific categorical (resp., continuous) feature. In the example, the weight 0.11 has the following meaning [27]: if the duration in months had been higher than the value it is for x , the prediction would have been, on average, 0.11 less “0” (or 0.11 more “1”). A not very easy logic to follow when compared to a single decision rule which characterize the contextual conditions for the decision of the black box. Finally, another advantage of our

$r = ((\text{credit_amount} > 836, \text{housing} = \text{own}, \text{other_debtors} = \text{none}, \text{credit_history} = \text{critical account}) \rightarrow \text{decision} = 0)$
 $\Phi = \{ ((\text{credit_amount} \leq 836, \text{housing} = \text{own}, \text{other_debtors} = \text{none}, \text{credit_history} = \text{critical account}) \rightarrow \text{decision} = 1),$
 $((\text{credit_amount} > 836, \text{housing} = \text{own}, \text{other_debtors} = \text{none}, \text{credit_history} = \text{all paid back}) \rightarrow \text{decision} = 1) \}$




Figure 7: Explanations of LORE (top) and LIME (bottom).

notion of explanation consists of the set of counterfactual rules. **LIME** provides an indication only of where to look for a different decision (different categorical values or lower/higher continuous values). **LORE**'s counterfactual rules provide an exact and minimal context for the prediction of a different decision.

5.4.2 Rule-Based vs Stochastic Counterfactuals. Regarding the counterfactual part of local explanations, we compare **LORE** with the stochastic optimization approach recently proposed in [33]. The approach looks for an instance x' as close as possible to a given x , but for which the black box outputs a different prediction. The Manhattan distance weighted by the inverse median absolute deviation is used in the search of x' . The approach returns x' as result. **LORE** is more general, since it provides one or more counterfactual

Dataset	Method	nf	c -hit	cl -fidelity
german	lore	1.52 \pm 1.18	.7765 \pm .38	.5355 \pm .43
	opt	14.80 \pm 1.59	.3118 \pm .47	.2297 \pm .36
compas	lore	1.84 \pm 0.78	.8694 \pm .37	.8611 \pm .41
	opt	6.24 \pm 1.45	.8036 \pm .38	.7555 \pm .34

Table 4: LORE vs OPT: performance of counterfactual rules.

rules, that generalize patterns of instances such as x' , hence providing the user with more informative output. In some sense, the stochastic optimization approach is limited to the phase of neighborhood generation, with a single element in the neighborhood. In order to make a fair comparison with our approach, we have implemented the stochastic optimization as an alternative fitness function of the genetic neighborhood generation. Table 4 shows the performances of the two methods on nf (number of falsified conditions in counterfactual rules), c -hit (rate of agreement of black box and counterfactual decision for counterfactual instance), and cl -fidelity (f1-measure of agreement of black box and counterfactual decision for instances covered by the counterfactual rule). The first is a measure of complexity of counterfactual explanations: **LORE** returns shorter explanations than OPT. The last two are measures of fidelity on the neighborhood instances that satisfy the premise of the counterfactual rules. They are considerably higher for **LORE**, especially for the german dataset. Once again, the merit is in the ability of **LORE** to generate a neighborhood that covers a dense and compact area around the instance under analysis. As a consequence, changes suggested by the counterfactuals tend to be small, and decision trees built from the neighborhood are not large.

6 CONCLUSION

We have proposed **LORE** a local black box agnostic explanator based on logic rules. **LORE** builds an interpretable predictor for a given black box and instance to be explained. The local interpretable predictor, a decision tree, is trained on a dense set of artificial instances similar to the one to explain generated by a genetic algorithm. The local decision tree enables the extraction of a local explanation, consisting of a single rule and a set of counterfactuals expressing which feature of the instance and how should be modified to get a different outcome. An ample experimental evaluation of the proposed approach has demonstrated the effectiveness of the genetic neighborhood procedure that leads **LORE** to outperform the proposals in the state of the art.

LORE is defined to treat tabular data. Thus, an interesting future research direction is to make the method suitable for image and text data for example by applying a preprocessing on the data for tagging them by semantic concepts, that might be mapped to a tabular format. Moreover, it would be interesting to investigate the possibility to derive a global description of the black box bottom-up by composing the local explanations and minimizing the size (complexity) of the global description.

REFERENCES

- [1] R. Andrews, J. Diederich, and A. B. Tickle. Survey and critique of techniques for extracting rules from trained artificial neural networks. *Knowl.-Based Syst.*, 8(6):373–389, 1995.
- [2] M. G. Augasta and T. Kathirvalavakumar. Reverse engineering the neural networks for rule extraction in classification problems. *Neural Processing Letters*, 35(2):131–150, 2012.
- [3] T. Bäck, D. B. Fogel, and Z. Michalewicz. *Evolutionary computation 1: Basic algorithms and operators*, volume 1. CRC press, 2000.
- [4] S. Barocas and A. D. Selbst. Big data’s disparate impact. *California Law Review*, 104, 2016.
- [5] R. Berk, H. Heidari, S. Jabbari, M. Kearns, , and A. Roth. Fairness in criminal justice risk assessments: The state of the art. *arXiv preprint arXiv:1703.09207*, 2017.
- [6] M. Craven and J. W. Shavlik. Extracting tree-structured representations of trained networks. In *NIPS*, pages 24–30. MIT Press, 1995.
- [7] J. Derrac, S. García, and F. Herrera. A survey on evolutionary instance selection and generation. *Int. J. of Applied Metaheuristic Computing*, 1(1):60–92, 2010.
- [8] F. Doshi-Velez and B. Kim. Towards a rigorous science of interpretable machine learning. *arXiv preprint arXiv:1702.08608v2*, 2017.
- [9] F.-A. Fortin, F.-M. De Rainville, M.-A. Gardner, M. Parizeau, and C. Gagné. DEAP: Evolutionary algorithms made easy. *Journal of Machine Learning Research*, 13:2171–2175, 2012.
- [10] Y. Fu, X. Zhu, and B. Li. A survey on instance selection for active learning. *Knowl. Inf. Syst.*, 35(2):249–283, 2013.
- [11] G. Fung, S. Sandilya, and R. B. Rao. Rule extraction from linear support vector machines. In *KDD*, pages 32–40. ACM, 2005.
- [12] B. Goodman and S. R. Flaxman. EU regulations on algorithmic decision-making and a “right to explanation”. volume abs/1606.08813, 2016.
- [13] R. Guidotti, A. Monreale, F. Turini, D. Pedreschi, and F. Giannotti. A survey of methods for explaining black box models. *arXiv preprint arXiv:1802.01933*, 2018.
- [14] S. Hara and K. Hayashi. Making tree ensembles interpretable. *arXiv preprint arXiv:1606.05390*, 2016.
- [15] A. Henelius, K. Puolamäki, H. Boström, L. Asker, and P. Papapetrou. A peek into the black box: exploring classifiers by randomization. *Data mining and knowledge discovery*, 28(5-6):1503–1529, 2014.
- [16] J. H. Holland. *Adaptation in natural and artificial systems: an introductory analysis with applications to biology, control, and artificial intelligence*. MIT press, 1992.
- [17] U. Johansson, L. Niklasson, and R. König. Accuracy vs. comprehensibility in data mining models. In *Int. Conf. on Information Fusion*, pages 295–300, vol. 1, 2004.
- [18] E. J. Keogh, S. Lonardi, and C. A. Ratanamahatana. Towards parameter-free data mining. In *KDD*, pages 206–215. ACM, 2004.
- [19] R. Krishnan, G. Sivakumar, and P. Bhattacharya. Extracting decision trees from trained neural networks. *Pattern recognition*, 32(12), 1999.
- [20] H. Lakkaraju, S. H. Bach, and J. Leskovec. Interpretable decision sets: A joint framework for description and prediction. In *KDD*, pages 1675–1684. ACM, 2016.
- [21] Y. Lou, R. Caruana, and J. Gehrke. Intelligent models for classification and regression. In *KDD*, pages 150–158. ACM, 2012.
- [22] D. M. Malioutov, K. R. Varshney, A. Emad, and S. Dash. Learning interpretable classification rules with boolean compressed sensing. In *Transparent Data Mining for Big and Small Data*, pages 95–121. Springer, 2017.
- [23] B. McCane and M. Albert. Distance functions for categorical and mixed variables. *Pattern Recognition Letters*, 29(7):986–993, 2008.
- [24] H. Núñez, C. Angulo, and A. Català. Rule extraction from support vector machines. In *Esann*, pages 107–112, 2002.
- [25] J. A. Olvera-López, J. A. Carrasco-Ochoa, J. F. Martínez Trinidad, and J. Kittler. A review of instance selection methods. *Artif. Intell. Rev.*, 34(2):133–143, 2010.
- [26] D. Pedreschi, S. Ruggieri, and F. Turini. Discrimination-aware data mining. In *KDD*, pages 560–568. ACM, 2008.
- [27] M. T. Ribeiro, S. Singh, and C. Guestrin. “why should I trust you?”: Explaining the predictions of any classifier. In *KDD*, pages 1135–1144. ACM, 2016.
- [28] S. Ruggieri. Yadt: Yet another decision tree builder. In *Tools with Artificial Intelligence, ICTAI*, pages 260–265. IEEE, 2004.
- [29] H. F. Tan, G. Hooker, and M. T. Wells. Tree space prototypes: Another look at making tree ensembles interpretable. *arXiv preprint arXiv:1611.07115*, 2016.
- [30] P.-N. Tan, M. Steinbach, and V. Kumar. Introduction to data mining. 1st, 2005.
- [31] C. Tsai, W. Eberle, and C. Chu. Genetic algorithms in feature and instance selection. *Knowl.-Based Syst.*, 39:240–247, 2013.
- [32] S. Wachter, B. Mittelstadt, and L. Floridi. Why a right to explanation of automated decision-making does not exist in the general data protection regulation. *International Data Privacy Law*, 7(2):76–99, 2017.
- [33] S. Wachter, B. Mittelstadt, and C. Russell. Counterfactual explanations without opening the black box: Automated decisions and the GDPR. *Harvard Journal of Law & Technology, Forthcoming*, 2017.
- [34] F. Wang and C. Rudin. Falling rule lists. In *AISTATS*, volume 38 of *JMLR Workshop and Conference Proceedings*. JMLR.org, 2015.
- [35] D. R. Wilson and T. R. Martinez. Improved heterogeneous distance functions. *J. Artif. Intell. Res.*, 6:1–34, 1997.
- [36] S. Wu. *Better Decision Tree from Intelligent Instance Selection*. VDM Verlag, 2009.
- [37] K. Xu et al. Show, attend and tell: Neural image caption generation with visual attention. In *ICML*, pages 2048–2057, 2015.
- [38] B. Zhou, A. Khosla, A. Lapedriza, A. Oliva, and A. Torralba. Learning deep features for discriminative localization. In *CVPR*, pages 2921–2929. IEEE, 2016.

Distance	hit	fidelity	l-fidelity	c	p	nf
cosine	.938±.24	.976±.11	.936±.24	4.4±2.5	2.1±1.8	1.9±1.0
minmax	.958±.19	.965±.15	.956±.17	4.5±2.7	2.3±2.3	1.8±0.9
neucld	.966±.17	.967±.15	.963±.19	4.3±2.6	2.2±2.1	1.8±1.0

Table 5: Comparison of distance measures on the *german* dataset.

A ADDITIONAL EXPERIMENTS

This section includes some further experimental results.

A.1 Comparing Distance Functions

A key element of the neighborhood generation is the distance function used by the genetic algorithm. A legitimate question is whether the results of the approach are affected by the choice of the distance. For instance, [33] reports considerable differences in their output of counterfactual instance based on the choice of the distance in their stochastic optimization approach. Table 5 reports basic measures contrasting the *normalized Euclidean* distance adopted by **LORE** with *cosine* and *min-max* distance. The table does not highlight any considerable difference. This can be justified by the fact that, following instance generation, there are phases, such as decision tree building, that abstracts instances to patterns, resulting in resilience against variability due to the distance function adopted.

A.2 Coverage of Global Features

In this section we replicate the experiment in [27] to measure the *faithfulness* of explanations with respect to classifiers that are interpretable by design, i.e., decision tree (DT) and logistic regression with L2 regularization (LR). However, differently from [27] we do not have to specify a fixed number of “gold” features, i.e., the most important features that can be used by the global model. To make a proper comparison, we rank the features by importance values returned by DT and LR. Then, we consider an increasing number of gold features going from the top-2 to the top-10, and use them in **LORE**. In our experiments, we study both recall over the black box, as in the global method of [27], and the precision.

Figure 8 reports the results of these analysis at the variation of the number of features included in data, for the *german* (first and second row) and *compass* (third and last row) datasets. Since LIME requires the number of gold features for the explanation, we run two versions of it: LIME which uses as number of features the same number of the gold features (from 2 to 10), and LIME-4 which uses 4 as number of features, corresponding to the average length of the premises of LIME. LIME, which is based on logistic regression, has a better performance with LR compared to DT. On the other hand **LORE**, which is based on decision trees, has better performance with DT compared to LR. Contrasting **LORE** and LIME, the former has better performance in terms of precision (due to its local nature), while the latter has better recall if 6 or more gold features are provided by the user (due to its global nature). **LORE** overcomes LIME in identifying the minimum set of gold features that should characterize an explanation, while LIME is better than **LORE** in recalling in an explanation the global features used by the black box when a considerable number of gold features is considered (and suggested to LIME).

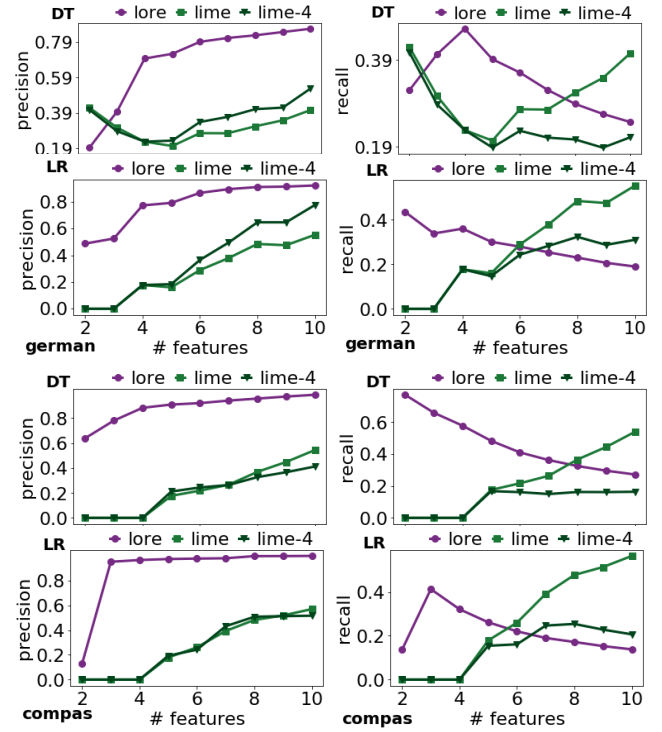


Figure 8: Precision and recall at the variation of the number of features.